

# TF-IDF and SVC Text Classification: *Is this text a song?*

**Aditya Mehta**  
A16062688  
a7mehta@ucsd.edu

**Isabel Suizo**  
A15485717  
isuiizo@ucsd.edu

**Emily Zhuang**  
A15592650  
ezhuang@ucsd.edu

## 1 Introduction

This report has been shortened to better fit the assignment requirements; please feel free to find an unabridged version of our paper [\[here\]](#).

## 2 Dataset

### 2.1 Collection

For this experiment, we constructed our own dataset to better suit the binary song lyric classification task at hand. The positive set of song lyrics was scraped from the popular song lyric website, Genius. First, we obtained a list of 160,000 songs from a Kaggle dataset which consists of songs released between 1921-2020, used that list to query for each song using the Genius API, and extracted the lyrics from each request. (Note: No data was used from the Kaggle dataset, as it was only used as a master list for songs to query.)

For the set of non-lyrics, we compiled data from several sources. A majority of the negative samples were taken from submissions on Reddit, the online discussion forum. We leveraged the variation of text content from different Subreddits to compile a subset of diverse text including poetry, personal stories, and professional writing (Table 4). We scraped the 27 Subreddits using the Reddit API to retrieve the most recent submissions and extracted the text data from each request. To vary the negative dataset even further, we also included reviews from the book database site Goodreads and the video game distribution platform Steam. Both review datasets were retrieved from Professor Julian McAuley’s database. Finally, we included shuffled instances of lyrics from the positive set, as well as a shuffled Reddit submissions and reviews.

While building our dataset, we made several intentional design choices throughout the process. First, we chose the Kaggle set as the master song

Table 1: Distribution of Negative Samples

Type	Count	% of Negative Set
Reddit	25000	64.1%
Goodreads	5000	12.8%
Steam	2000	5.1%
Shuffled Song	4000	10.3%
Shuffled Reddit	1000	2.6%
Shuffled Goodreads	1000	2.6%
Shuffled Steam	1000	2.6%

list of positive lyric samples since it was a very comprehensive list that featured a reasonable spread of genres across the past century until 2020. Because of time constraints, we decided to select only a subset of song lyrics from the master list. To maintain an even spread of samples, we first shuffled the master list before we extracted 40,000 samples. For our non-song lyrics, we reasoned that pulling from various Subreddits would generate a diverse collection of text genres since we chose Subreddits that each had a different theme of varying lengths. A small subset of our negative set also consists of reviews from Goodreads and Steam. These samples were added for more diversity in text content since our Reddit submissions do not include review-style text. Furthermore, our negative sample also includes shuffled song lyrics. This set was crucial to training our dataset to encourage our model to also take sequence into consideration. Therefore, we included this shuffled set to motivate extra attentiveness to the structure of song lyrics.

### 2.2 Pre-Processing

Once we had collected our raw data, we had quickly noticed that we would have to consider multiple things. How would we handle non-English samples, emojis, internet language, and

Table 2: Train/Test/Validation Splits After Preprocessing. The last column represents the percentage of samples from the overall dataset.

Set	Count (Lyrics/Non-Lyrics)	%
Training	31330/30744	80%
Validation	3916/3843	10%
Test	3916/3843	10%

certain tags that are characteristic to only the song lyrics set?

To try to reduce the number of non-English samples from our dataset, we settled on using a 30% threshold to determine whether a sample would be discarded from the dataset. If 30% or more of the sample were words that did not exist in the English dictionary, we threw the sample out. Since our raw song lyric data was predominantly English song lyrics, we wanted to limit the scope of our predictive task to focus on English text samples. The reason why we had to use a threshold as opposed to just throwing out all samples that contained a single non-English word was because many song lyrics include slang words or shortened forms of words that are not technically English words, but we still wanted to keep those positive samples in our dataset since that is characteristic of many songs and it would be oversimplifying the predictive task. It is also characteristic of many other forms of English written text to include variations of words, acronyms, initialisms, slang, and, in the recent years, emojis that we, as a society, have collectively assigned meaning to such groupings of letter. Therefore, to create a generalized model that would be able to distinguish between song lyrics and non-song lyric text, we wanted to ensure our dataset included enough of these non-English words to be functional.

We also noticed many, but certainly not all, of the song lyric samples included tags within the text like, "[Chorus]" and "[Verse 1]", to indicate sections of the song. It is obvious that the majority of the time these tags appear, the sample is a song's lyrics. Therefore, the question we had to ask ourselves was, do we remove said tags since it might make the model a trivial predictor? More precisely, we were worried that leaving the tags in may cause the model to look only for such tags within the text and predict that a chunk of text came from song lyrics if and only if it saw one of those tags. In the end, we decided to leave those

tags in for the time being to see how the model performed with the tags left in (SPOILER ALERT: we never ended up removing the tags from the positive sample set). Since we intended to build a robust model, the model should be able to handle these complexities. Also, since it is very common for song lyrics to contain tags, it might in fact be a good indicator to the model that a set of text is a song's lyrics if it contains tags.

The last pre-processing we did was to randomly shuffle the words a small subset of presamples from the data that we had obtained from each of the other sources. We took 4,000 song lyric samples, 1,000 Goodreads reviews, 1,000 Reddit posts, and 1,000 Steam reviews to build some more negative samples that would give our model some more complexities since we wanted to train it to identify the difference between the original song lyrics and an un-ordered collection of the same words that appear in song lyrics. That is also why we included a larger sample of song lyrics within our set of shuffled samples. We were thinking that the model would have more of a difficulty differentiating between the scrambled and unscrambled song lyrics than differentiating between scrambled negative samples and unscrambled song lyrics, so we wanted to give the model more scrambled song lyric samples to work with.

## 2.3 Exploration

Now that we have collected and cleaned our dataset, it is time to unpack the features of our data. To have a comprehensive understanding of our data looks, we visualize specific characteristics of the text data using various tables and graphs. We will explore the most common words and bigrams that appear, the most common words that appear in only the non-song lyric text and not song lyrics, and the distribution of text length within the dataset.

To better understand the difference between lyrics and non-lyrics, we analyzed the most common words that appear in non-song lyric text, but do not appear in song lyrics as noted in Table 5. First, several acronyms are shown in Table 5 such as "lmao" and "tldr". The absence of these acronyms in lyrics can be attributed to the fact that they are not true words with phonemes, hence would not fit in a melodic setting. Additionally, our negative set consists entirely of text examples scraped from websites where use of

acronyms is very common. Furthermore, we also observe several names in this dataset, including "mallory", "ashlyn", "becca". We presume this is because our negative consists of several storytelling Subreddits which could reference specific names. Also, the reviews from Goodreads and Steam may also feature discussion on characters of novels or videogames. Finally, we can also see the use of modern terms like "ghosting" and "red-ditors." Since these words have grown in popularity by propagating through the internet, it makes sense that these words would show up in our non-lyric text dominated by internet content. We also note the frequency of binary sequences in our non-lyric text. These sequences represent characters which are frequently seen in internet contexts.

We also analyzed the differences between the most common unigrams and bigrams in each half of the dataset. Figure 6 and Figure 7 display the most common 100 words within song lyrics and non-song lyric text and the words' overall frequencies within each sample set. The majority of the words within both common word sets are stop words and the overlap between the sets is large. Meaning, within the most common words, there are not many distinct words that are characteristic of either half of the dataset. Words like "love", "time", and "life" are all words that appear very frequently within song lyrics but not as frequently in non-song lyric text. Figure 8 and Figure 9, showing the 100 most common bigrams within song lyrics and non-song lyric text, don't contribute much more insight to the nature of the dataset. The high frequency of stop words and pronouns in unigrams persists to bigrams. Figure 10 and Figure 11, the 101th-200th most common bigrams in each dataset, offer a better idea of the contents of each text set. We can see that song lyrics frequently contain "i love", "love you", "my heart", and "in love" which makes sense since a large subset of songs are about love, while the non-song lyric text contain a lot of "she had", "he said", "she said", "with her", and "as he" since many of the Reddit samples discuss the interactions between individuals. As a generalization, we can conclude that song lyrics tend to discuss abstract ideas like emotions and ideas while non-song lyric text more frequently talks about concrete actions and items.

We also visualized the distribution of the length of each body of text as depicted in Figure 1 to ver-




Figure 1: Distribution of Logarithmic Transformation of Text Length (in characters) Across Sets

ify the distribution of our personalized non-song lyric text dataset. As mentioned in the dataset section, we wanted to collect text samples of varying length that resemble the distribution of our song lyrics. Based on Figure 1, we observe that the median of the negative set is only slightly below the median of the lyrics. The negative review set shows a greater distribution in length, however, this can be explained by the variation in the length of our different non-lyric sources. While the Steam and Goodreads reviews were relatively short, some Subreddits like nosleep had an average length of 12453 characters which is well above the song lyric average of 5466 characters.

### 3 Predictive Task

While all the data we collected is consistent in that each entry is a large block of text, it is also separated into two distinct parts – song lyrics and non-song lyric text. We decided to center our predictive task around the binary nature of our dataset by focusing on a model which would predict whether or not large blocks of text are song lyrics.

#### 3.1 Baseline

Our initial intuition was to determine if a block of text identifies as song lyrics based on whether or not it contains words specific to either classification. If a block of text contains a word(s) that appears exclusively in non-song lyric texts, it might be reasonable to predict it as a non-song lyric block of text.

We first collected a set of unique words that appeared in all of the song lyrics in the training set

and all of the non-song texts in the training set. We then took the difference of the non-song set from the song set, which produced a set of words which appeared in the non-song set but *did not* appear in the song set.

To predict, we iterated through each entry in the training set and predicted it to be a non-song if even a single word from the exclusive non-song set appeared in the text block. Otherwise, we predicted it to be a song. This resulted in an accuracy of 65.64299424184261%. Above average, but clearly sub-optimal, this score provided a robust starting point to build a complex model off of.

### 3.2 Evaluation Metrics

Our evaluation metric will be prediction accuracy and BER for judging the completeness of our task. We agreed that the more often it predicts correct labels, the better the model. Thus, we settled on finding the best model through the highest accuracy score when using the equation:

$$Score = \frac{|yPredictions^{TP}| + |yPredictions^{TN}|}{|yPredictions|}$$

We also used balanced error rate as an additional metric. Since BER is an average of false positive and false negative rates, it is more telling of success in the case of imbalanced data sets. In our case, accuracy is a completely reasonable metric to judge our model given our datasets each have roughly the same number of positive and negative samples. However, we decided to also incorporate BER in case it captured slight variations that accuracy did not reveal.

### 3.3 Feature Extraction and Model Choice

When choosing a model, we knew that choosing a text-feature based design was an absolute requirement. Because of our text-only dataset, we had to create a model which would be able to push out strong predictions based on text features alone. The naive implementation of this idea, creating a bag of words vector, is slow, inefficient, and was never considered. We recognized the importance of bigrams and trigrams in our text, as groupings of words could go a long way in song recognition. Using TF-IDF as our feature vectorizer would not only allow us to take advantage of a strong text-feature extractor, it would also allow us to fine tune important feature settings such as dictionary size and n-gram range.

Next, we considered which mathematical model to implement as our predictor. Logistic regression seemed like a good fit for our task because of its binary nature, which mirrored our binary classifications. Using a support vector classification also had its strengths - namely the built-in associated learning algorithms which would help us fine tune our accuracy. After weighing both options, we decided to experiment with both predictors before finalizing a model.

Finally, we considered means of extracting features from large blocks of text. Text tokenizations such as removing stop words, removing punctuation, stemming, removing emojis, and converting to all lowercase were all viable approaches which could help the TF-IDF vectorizer extract the most relevant features. We planned to experiment with our model's feature parameters so that it could consider different dictionary sizes (ranging from 5000 to 25000) as well as combinations of n-gram ranges (lower/upper bounds ranging from 1-3).

## 4 Model

We decided to implement several different models before fixating on a singular optimal model. We divided the task of finalizing a model into several stages: finding the best feature settings, tuning hyperparameters on a model with these best feature settings, and finally combining the tuned model with SVC and Logistic Regression. Although we decided to keep many different variations of this model in play, one factor was consistent throughout - using a TF-IDF Vectorizer. The decision to use the TF-IDF Vectorizer can be easily justified because of how powerful it is in extracting key details from blocks of text, all while considering different n-gram ranges and dictionary sizes.

### 4.1 Architecture

Our model utilizes several different approaches to ensure optimization. Firstly, the decision to use TF-IDF is already a huge optimization over other popular approaches to text feature extraction. Unlike a naive implementation which may use a bag-of-words approach to extract features, TF-IDF is able to create feature vectors from blocks of text with more efficiency. TF-IDF is able to prioritize and weigh words and compact the already smaller feature vectors into sparse-matrices. Its ability to do this greatly improves the efficiency and speed of the model as opposed to other approaches.

Secondly, because of our extensive dataset, we were able to train the model on large volumes of text, with the training set containing 63,206 entries. We ensure that the split of song lyrics and non-song text within these 63,206 entries is nearly even, which means the training on the model is extremely balanced.

Additionally, we will optimize our model(s) further in the second stage of our model-selection process. As described above, this will revolve around tuning the hyperparameters on the model after its best feature settings have been discovered. This is crucial because we will be specifically experimenting with regularization to increase the model's resistance to overfitting, as well as other types of fine tuning.

Some issues we ran into when considering different models predominantly revolved around scalability. We found that when trying to create a predictor which considered quadgrams (n-grams with a range up to 4 words), the computer's RAM would overload and the model would crash. This would also happen when trying to create a feature vector with over 40,000 entries. Crashes like these were definitely expected as we tried to push the boundaries of the model's scalability. While more features and n-grams definitely increased the scope of the model's learning, it became quickly apparent that there was a limit to this scope.

## 4.2 Model Trials

Our baseline model proved to be our most unsuccessful attempt; not only did the model have sub-optimal accuracy, it also had a high BER of over 34%. It was clear that we had reached the ceiling in terms of performance when using this approach. The baseline's approach, which predicted labels whether or not the test set text contained words exclusive to non-song text, relied too heavily on the words that showed up in the training set. Additionally, it also required extreme similarity between the train and test set in order to work, such that if the test set contained no words that appeared in the training set, the model would predict 'song lyric' for every single entry. After realizing that the baseline could no longer be optimized further, we pivoted to a different type of text-feature extraction which still retained the overall idea of text analysis through word composition – TF-IDF.

We decided to consider several different types of models for the model that had the best fea-

ture setting and tuned hyperparameters. We decided to compare this model using several predictors such as the baseline, logistic regression and support vector classification.

For the baseline model, the weaknesses definitely outweighed the strengths (as discussed above). However, it was a valuable model for our process because it gave us a foundation on how we wished to approach this predictive task – by analyzing word patterns within the text. When considering the support vector classifier, a key strength that stood out was its ability to be more generalized in its prediction pattern, which would ensure resistance to overfitting. On the other hand, there was an argument to be made against SVC's, as they take longer to run compared to predictors that use linear regression. This could prove to be an issue for us in terms of scalability, as we are dealing with a large dataset to train on. Finally, using logistic regression for our predictor had appeal because of how straight-forward it is. Unfortunately, its speed and readability are opposed by weaknesses such as a lack of the precision and associated learning algorithms that support vector classifiers provide.

After considering these factors, we decided that our final model would be one which identified the best feature settings, fine tuned the hyperparameters maximally, and chose the best predictor between a linear regressor and support vector classifier.

## 5 Literature

The nature of our predictive task necessitated that we conduct research about where and how we would obtain our dataset, how would we study our data in a meaningful way, and which cutting edge methods have been shown to perform well on similar data and for similar tasks. We will break down all of our references in this section that will explain much of our model's design.

### 5.1 Existing Dataset

As detailed in our Data Collection subsection, we compiled our own dataset. One of the existing datasets we did use, a Kaggle dataset, was used only as a comprehensive list of songs to use for our query of song lyrics [ay]. The original Kaggle dataset included song data like artists, song duration, release year, danceability, acousticness, tempo, popularity, and song key. Many other in-

dividuals used this data for other predictive tasks like predicting a song's popularity, genre, or release year. This paper also uses two datasets from Professor McAuley's database, each containing review data from Goodreads ([Kang and McAuley, 2018], [Pathak et al., 2017],[Wan and McAuley, 2018]) and Steam [Wan and McAuley, 2018]. Both datasets were used to extract review text to add greater variation to our non-lyric samples.

## 5.2 Similar Datasets

Other similar datasets have compiled song lyrics to be used for song lyric generation. This other Kaggle dataset, grouping the works of many songs by artist and poems by poet, has been used to generate song lyrics for specific artists [Mooney, 2018]. For example, Amisha Jodhani used the dataset to generate songs that mimic the characteristics of Bruno Mars's songs using a neural network [Jodhani, 2020].

Another individual also analyzed a song lyric dataset in nearly way, shape, and form and wrote about his/her findings in a Kaggle discussion post [Kratisaxena, 2018]. He/she details the distributions of number of songs by artist, the song lyric length, and song title length. He/she even dives deeper and explores the relationship between song title length and song length, the words used to express different sentiments within songs, the most common words used by a specific artist, and also any common rhythmic words. Although the author of this piece explores all aspects of song lyrics, the techniques he/she uses to analyze the data are fairly standard. The use of frequency tables and n-grams is fairly consistent with other studies of text datasets.

## 5.3 Prior Research

For this binary text classification problem, our results and research both pointed to the success of Support Vector Classifiers (SVC). According to this paper, this success in high-dimensional spaces is explained by their ability to learn "independent of the dimensionality of the features" [Joachims, 1998] by biasing towards a margin that can partition our data. This quality is perfect for text classification given our feature representations are usually thousands of words long where each feature corresponds to some important word in our corpus.

## 5.4 State-of-the-Art Methods

In recent years, we've observed the exponential growth in the implementation of neural networks for these text categorization problems. According to PapersWithCode, the best model in 2019 for text categorization of question topics on the website, Yahoo! Answers, leverages a pretrained BERT to achieve 77.62% accuracy [pap, 2019]. Typically BERT, short for bidirectional encoder from transformers, is used for textual entailment tasks to predict the next word in a sequence, but this paper was able to adapt BERT to solve a categorization task [Sun et al., 2019]. This model was able to achieve state-of-the-art results and reduce the error of modern solutions by 18.57% on average through utilizing pretrained WordPiece embeddings and identifying the layer of most effective features.

Beyond neural networks, we have also seen attempts to solve this categorization problem through hierarchical classification. According to the paper, "Hierarchical Text Classification and Evaluation," modern solutions only consider categories of text as isolated entities, as opposed to a natural hierarchy where a document can be classified into subcategories down the tree until it reaches a leaf category [Sun and Lim, 2001]. In order to classify a document, this method has binary classifiers at each category to determine whether the document should belong to the corresponding category, or if it should be given to a classifier in its subcategories. This method proved to be reasonably successful based on the personalized error metrics. However, the paper also suggests this novel solution has much room for improvement since it does not take full advantage of hierarchical properties and suffers from irrecoverable errors made by parent-category classifications.

## 5.5 Similar Papers

A Stanford paper, "Using Song Lyrics and Frequency to Predict Genre" explores a similar task of categorizing rap vs. non-rap lyrics [Ram, 2017]. To generate its feature vector, this method simply used 500 words that had a frequency of 1000 to 3000 occurrences in the training set and counted the frequency of each word for each sample. This experiment achieved similar results with 93.2% accuracy in rap lyric classification using SVM, also outperforming their naive bayes and regression models. Our model slightly outperforms their



Feature Variations	N-gram Range	Dictionary Size	Test Accuracy
<i>original (no extra processing of inputs)</i>	[1, 1]	5000	0.9151357279956128
<i>remove stop words</i>	[1, 1]	5000	0.9145873320537428
<i>remove punctuation</i>	[1, 1]	5000	0.9158212229229503
<i>stemming</i>	[1, 1]	5000	0.9141760350973402
<i>remove emojis</i>	[1, 1]	5000	0.9152728269810804
<i>convert to lowercase</i>	[1, 1]	5000	0.9151357279956128
<i>remove stop words, punctuation, and emojis, stemming, and convert to lowercase</i>	[1, 1]	5000	0.9123937482862626

Figure 2: Test Accuracies of SVC Model with Various Input Processing

model, which can be attributed to our use of TF-IDF vectorization as well as the comparative difficulty of our classification task.

## 6 Results

We found that the model with the best test accuracy was the model using stemming on the input text, unigrams, bigrams, and trigrams and a feature size of 25,000 for the TF-IDF Vectorization, and the feeding the feature vector generated by the TF-IDF into an SVC model.

### 6.1 The Numbers

We identified that the best performing model for our dataset and specific predictive task was an SVC model that used TF-IDF with n-gram range [1, 3], feature size 25,000, and stemming the input text to build the feature vector. This model was able to differentiate song lyrics from non-song lyric text with an accuracy of 0.97203 and a BER of 0.02810. With the best performing model highlighted in green, Figure 2, Figure 3, Figure 4, and Figure 5 show the results of various combinations of different models, input processing, and feature design. The figures only show a subset of the results we obtained. Since we wanted to obtain the highest performing combination of different feature optimizations, we performed a grid search over 84 different models while making single variable changes to n-gram range, dictionary size, and input text processing in each iteration. We have condensed the results of the grid search to focus on the key trends.

### 6.2 Feature Reflection

Some notable trends within our results is that overall, a larger range of n-grams and a larger dictio-

Feature Variations	N-gram Range	Dictionary Size	Test Accuracy
<i>original (no extra processing of inputs)</i>	[1, 3]	25000	0.9714834110227585
<i>remove stop words</i>	[1, 3]	25000	0.9317247052371812
<i>remove punctuation</i>	[1, 3]	25000	0.9713463120372909
<i>stemming</i>	[1, 3]	25000	0.9720318069646284
<i>remove emojis</i>	[1, 3]	25000	0.9716205100082259
<i>convert to lowercase</i>	[1, 3]	25000	0.9714834110227585
<i>remove stop words, punctuation, and emojis, stemming, and convert to lowercase</i>	[1, 3]	25000	0.9422813271181794

Figure 3: Test Accuracies of SVC Model with Various Input Processing

Model Type	Features Type	N-gram Range	Dictionary Size	Test Accuracy
<i>SVC</i>	stemming	[1, 1]	5000	0.9141760350973402
<i>SVC</i>	stemming	[1, 1]	25000	0.9159583219084179
<i>SVC</i>	stemming	[1, 2]	5000	0.9621606800109679
<i>SVC</i>	stemming	[1, 2]	25000	0.9707979160954209
<i>SVC</i>	stemming	[1, 3]	5000	0.9636687688511105
<i>SVC</i>	stemming	[1, 3]	25000	0.9720318069646284
<i>SVC</i>	stemming	[2, 2]	5000	0.9439265149437894
<i>SVC</i>	stemming	[2, 2]	25000	0.9642171647929806
<i>SVC</i>	stemming	[2, 3]	5000	0.9433781190019194
<i>SVC</i>	stemming	[2, 3]	25000	0.9628461749383055
<i>SVC</i>	stemming	[3, 3]	5000	0.8925143953934741
<i>SVC</i>	stemming	[3, 3]	25000	0.9284343295859611

Figure 4: Test Accuracies of SVC Model with Various Features

Model Type	Feature Type	N-gram Range	Dictionary Size	Test Accuracy	Test BER
<i>baseline</i>	positive negative set comparison	N/A	N/A	0.6564299424184261	0.34777122618186473
<i>logistic regression</i>	stemming	[1, 3]	25000	0.9606525911708254	0.03956517119172476
<i>SVC</i>	stemming	[1, 3]	25000	0.972031806964628	0.028096151715940598
<i>SVC</i>	stemming and newline character count	[1, 3]	25000	0.9709350150808884	0.02923785351338759

Figure 5: Test Accuracies and Test BERs of Various Models Holding Features Constant

nary size performed better than unigrams and a small dictionary size. This is significant because, as intuition would tell us, song lyrics are highly structured forms of text that rely on specific repetitions of sounds and meter to create the lyrical flow that songs contain. This order is not maintained if we only inspect the contents of each input text using unigrams. With unigrams, the model would not be able to differentiate between song lyrics and the scrambled versions of song lyrics. A large dictionary size is also important because we want to include n-grams that potentially could appear in the dataset at a low frequency but could also be highly representative of either a song lyric or non-song lyric text as we had explored in the data exploration section. Also, since we are using a range of n-grams, this increases the number of unique values in our dictionary. Naturally, the size of our dictionary must increase to accommodate the complexities of the data.

We also found that out of the various different methods we used to process the input text, stemming the text performed marginally better than other variations while holding all other feature parameters constant. This is probably an effect of the fact that stemming is able to compress the information within the text and therefore allowed us to pass a more dense, informative feature vector to the model. Stemming text will group different words with the same stem as one value in the dictionary.

Removing stop words worked the least well out of all the variants by far when the n-gram range was [1, 3] and the dictionary size was 25,000. If we look at Figure 6 and Figure 7, we can see that the stop words appear in song lyrics at a significantly higher frequency than non-song lyric text. The frequency at which stop words occur in song lyrics is an order of magnitude greater than the frequency in non-song lyric text. This characteristic of the dataset is likely captured by the n-grams and removing stop words is removing a meaning indicator that the model uses to make its prediction.

Our models all employed TF-IDF to build the feature vector. After some thorough research and experimentation, we decided to use TF-IDF to build our feature vector because it would both allow us to leverage the vast array of words that we discovered appear in the non-song lyric dataset at relatively low frequencies as well as allow us to make our model resistant to a large range of input

Table 3: Accuracy of hyperparameter experiments

	<b>Regularization (C)</b>			
<b>Loss Function</b>	0.01	0.1	1	10
hinge	0.898	0.956	0.973	0.970
squared hinge	0.931	0.968	0.974	0.971

text length. The TF-IDF models performed nicely.

### 6.3 Model Parameters

To tune our hyperparameters, we ran a 2-dimensional grid search on various combinations of regularization constants and loss functions. We found that the regularization parameter  $C=1$  performs the best on the validation set. Tuning the regularization parameter is simply a tradeoff in weighing model complexity vs. actual prediction errors in our loss function. Therefore, we found the perfect sweet spot of just enough regularization without a major compromise on the model's classification ability.

In each of our experiments, squared hinge loss also performed better than hinge loss. This can be explained by the mechanics of squared loss which penalizes predictions with respect to their distance from the target label. This allows the model to converge to a solution that is sensitive to outliers, which would be highly applicable for our dataset with a large variation in textual features.

### 6.4 Comparative Analysis

Based on our results, SVC performed better than using a Logistic Regression classifier. This comes as no surprise since SVC is an implementation of a Support Vector Machine(SVM) which is known to perform better for this particular classification task. SVMs simply optimize a margin to partition samples of different classes. Since SVM uses this geometrical approach, as opposed to Logistic Regression which utilizes a probabilistic approach, it performs much better in high-dimensional spaces, such as our sample text feature space. In addition, SVMs are known for generalizing better since they do not penalize correct predictions made with reasonable confidence due to their sole objective of partitioning our categories with the greatest margin. In contrast, Logistic Regression uses a cross-entropy loss function where correct predictions still incur a cost. In this case, the characteristics of our features, in this case, the high-dimensional, sparse feature space, are most suitable for an SVM



classifier.

## References

Spotify dataset 1921-2020, 160k+ tracks.

2019. [Papers with code - yahoo! answers benchmark \(text classification\)](#).

Thorsten Joachims. 1998. [Text categorization with support vector machines: Learning with many relevant features](#). *Machine Learning: ECML-98 Lecture Notes in Computer Science*, page 137–142.

Amisha Jodhani. 2020. [Bruno mars song generator](#).

Wang-Cheng Kang and Julian McAuley. 2018. [Self-attentive sequential recommendation](#).

Kratisaxena. 2018. [Let’s analyze the song lyrics!](#)

Paul Mooney. 2018. [Song lyrics](#).

Apurva Pathak, Kshitiz Gupta, and Julian McAuley. 2017. [Generating and personalizing bundle recommendations on steam](#).

Jayen Ram. 2017. [Using song lyrics and frequency to predict genre](#).

Aixin Sun and Ee-Peng Lim. 2001. [Hierarchical text classification and evaluation](#). *Proceedings 2001 IEEE International Conference on Data Mining*.

Chi Sun, Xipeng Qiu, Yige Xu, and Xuanjing Huang. 2019. [How to fine-tune bert for text classification?](#) *Lecture Notes in Computer Science Chinese Computational Linguistics*, page 194–206.

Mengting Wan and Julian McAuley. 2018. [Item recommendation on monotonic behavior chains](#).

Table 5: Most Common Words That Appear in the Negative Set But Not the Positive Set

Table 4: Subreddits Scraped to Generate Negative Set. The count indicated the number of submissions used in our dataset and average length represents the average character count of each Reddit submission.




Figure 6: Number of Instances Words Appear in Song Lyrics




Figure 7: Number of Instances Words Appear in Non-Song Lyric Text




Figure 8: Number of Instances Bigrams Appear in Song Lyrics




Figure 9: Number of Instances Bigrams Appear in Non-Song Lyric Text




Figure 10: Distribution of Newline Characters in Non-Song Lyric Text




Figure 11: Distribution of Newline Characters in Song Lyrics





Figure 12: Frequency of 101th-200th Most Popular Bi-grams in Non-Song Lyric Text

Figure 13: Frequency of 101th-200th Most Popular Bi-grams in Song Lyrics